

Problem 15: Monte Carlo estimation of an expected value

Show that $\mathbb{E}(\hat{g}(X)) = \mathbb{E}(g(X))$:

$$\begin{aligned}\mathbb{E}(\hat{g}(X)) &= \mathbb{E}\left(\frac{1}{N} \sum_{i=1}^N g(X_i)\right) \\ &= \frac{1}{N} \sum_{i=1}^N \mathbb{E}(g(X_i)), \text{ by linearity of} \\ &\text{expectation. Since all } X_i \text{ are drawn}\end{aligned}$$

from the distribution X ,

$$\mathbb{E}(g(X_i)) = \mathbb{E}(g(X))$$

So,

$$\begin{aligned}&= \frac{1}{N} \sum_{i=1}^N \mathbb{E}(g(X)) = \frac{1}{N} \cdot N \mathbb{E}(g(X)) \\ &= \mathbb{E}(g(X))\end{aligned}$$

Show that $\text{Var}(\hat{g}(X)) = \frac{\text{Var}(g(X))}{N}$

$$\begin{aligned}\text{Var}(\hat{g}(X)) &= \text{Var}\left(\frac{1}{N} \sum_{i=1}^N g(X_i)\right) \\ &= \frac{1}{N^2} \text{Var}\left(\sum_{i=1}^N g(X_i)\right), \text{ since } 1/N \text{ is} \\ &\text{constant. By Bienaymé's identity:}\end{aligned}$$

$$\begin{aligned}&= \frac{1}{N^2} \text{Var}\left(\sum_{i=1}^N g(X_i)\right) = \frac{1}{N^2} \sum_{i=1}^N \text{Var}(g(X_i)) \\ &= \frac{1}{N^2} \sum_{i=1}^N \text{Var}(g(X)) = \frac{1}{N^2} \cdot N \cdot \text{Var}(g(X)) \\ &= \text{Var}(g(X)) / N\end{aligned}$$

ca)

$$P(C=T | R=T, S=T, W=T)$$

$$= P(C=T | S=T, R=T), \text{ by Markov blanket}$$

$$= \frac{P(C=T) P(S=T, R=T | C=T)}{P(S=T, R=T)}$$

$$= \frac{P(C=T) P(S=T | C=T) P(R=T | C=T)}{P(C=T) P(S=T | C=T) P(R=T | C=T) + P(C=F) P(S=T | C=F) P(R=T | C=F)} \quad (*)$$

Plug in values:

$$= \frac{0.5 \cdot 0.1 \cdot 0.8}{0.5 \cdot 0.1 \cdot 0.8 + 0.5 \cdot 0.5 \cdot 0.2}$$

$$= 0.4444 \dots$$

For $P(C=T | R=F, S=T, W=T)$ is completely analogous. Take (*) and replace $R=T$ with

$$R=F:$$

$$P(C=T | R=F, S=T, W=T)$$

$$= \frac{P(C=T) P(S=T | C=T) P(R=F | C=T)}{P(C=T) P(S=T | C=T) P(R=F) + P(C=F) P(S=T | C=F) P(R=F | C=F)}$$

$$= \frac{0.5 \cdot 0.1 \cdot 0.2}{0.5 \cdot 0.1 \cdot 0.2 + 0.5 \cdot 0.9 \cdot 0.8}$$

$$\approx 0.04762$$

We now compute

$$P(R=T | C=T, S=T, W=T)$$

$$= \frac{P(R=T | C=T, S=T) P(W=T | R=T, S=T)}{P(W=T | C=T, S=T)}$$

$$= \frac{P(R=T | C=T) P(W=T | R=T, S=T)}{P(R=T | C=T) P(W=T | R=T, S=T) + P(R=F | C=T) P(W=T | R=F, S=T)} \quad (C^{**})$$

Plug in values:

$$= \frac{0.8 \cdot 0.99}{0.8 \cdot 0.99 + 0.2 \cdot 0.9} \approx 0.8148$$

Analogously, we can find

$P(R=T \mid C=F, S=T, W=T)$ by replacing

$C=T$ with $C=F$ in (**)

$$= \frac{P(R=T \mid C=F) P(W=T \mid R=T, S=T)}{P(R=T \mid C=F) P(W=T \mid R=T, S=T) + P(R=F \mid C=F) P(W=T \mid R=F, S=T)}$$

$$= \frac{0.2 \cdot 0.99}{0.2 \cdot 0.99 + 0.8 \cdot 0.9} \approx 0.2157$$

Project 6 - Problem 16 (b) - (i)

Group J

2023-04-05

Subtask (b)

The probabilities derived in (a) are used to implement a Gibbs sampler for the network:

```
n <- 100
retMat <- matrix(ncol = 2, nrow = n)

#We define a matrix, where sampling a 0 or 1 corresponds to C or R respectively.
#The matrix contains the needed conditional probabilities to perform Gibbs sampling

probs <- matrix(data = c((.5*.1*.2)/(.5*.1*.2 + .5*.5*.8)),
                (.2*.99)/(.2*.99+.8*.9),
                (.5*.1*.8)/(.5*.1*.8 + .5*.5*.2),
                (.8*.99)/(.8*.99+.2*.9)), nrow = 2, ncol = 2)

#Define x_0 as the initial state of the Markov chain
x_0 <- c(1,1)
x <- x_0

#Uniformly sample whether to update C or R

#Get initial index
choice <- rbinom(1, 1, .5)
idx <- choice + 1

for (i in 1:n) {

  sampling_prob <- probs[idx, (!choice) + 1]
  sample <- rbinom(1, 1, sampling_prob)

  retMat[i, ] <- x
  x[idx] <- sample

  idx <- ifelse(idx == 2, yes = 1, no = 2)

}

colnames(retMat) <- c("C", "R")
head(retMat)
```

```
##      C R
```

```
## [1,] 1 1
## [2,] 1 0
## [3,] 0 0
## [4,] 0 0
## [5,] 0 0
## [6,] 0 0
```

Subtask (c)

This allows us to calculate the marginal probability $P(R = T|S = T, W = T)$:

```
counts <- matrix(data = c(0, 0, 0, 0),
                  nrow = 2, ncol = 2)

for (i in 1:n) {
  temp <- retMat[i, ]
  temp <- temp + 1
  counts[temp[1], temp[2]] <- counts[temp[1], temp[2]] + 1
}

# Calculate probability:
(task_c <- rowSums(counts/n)[2])
```


```
## [1] 0.06
```

Subtask (d)

Here we use the effective sample size (ESS) to estimate the autocorrelation in the data. Further, the **acf**-function plots the autocorrelation of each sample with its neighboring samples.

```
# We use the upper limit for k that is selected by the acf-function
acf1 <- acf(retMat[, 1])
```

Series retMat[, 1]




```
n/(1 + 2 *sum(acf1$acf))
```

```
## [1] 14.18226
```

```
acf2 <- acf(retMat[, 2])
```

Series retMat[, 2]



```
n/(1 + 2 * sum(acf2$acf))
```

```
## [1] 36.65333
```

Subtask (e)

```
sim_1 <- matrix(ncol = 2, nrow = 50000)
sim_2 <- matrix(ncol = 2, nrow = 50000)

#We define a matrix, where sampling a 0 or 1 corresponds to C or R respectively.
#The matrix contains the needed conditional probabilities to perform Gibbs sampling

probs <- matrix(data = c((.5*.1*.2)/(.5*.1*.2 + .5*.5*.8)),
                 (.2*.99)/(.2*.99+.8*.9),
                 (.5*.1*.8)/(.5*.1*.8 + .5*.5*.2),
                 (.8*.99)/(.8*.99+.2*.9)), nrow = 2, ncol = 2)

#Define x_0 as the initial state of the Markov chain
x_0 <- c(1,1)
x <- x_0

#Uniformly sample whether to update C or R
idx <- choice + 1
```



```

for (i in 1:50000) {

  sampling_prob <- probs[idx, (!choice) + 1]
  sample <- rbinom(1, 1, sampling_prob)

  sim_1[i, ] <- x
  x[idx] <- sample

  idx <- ifelse(idx == 2, yes = 1, no = 2)
}

x_0 <- c(1,1)
x <- x_0

#Uniformly sample whether to update C or R
idx <- choice + 1

for (i in 1:50000) {

  sampling_prob <- probs[idx, (!choice) + 1]
  sample <- rbinom(1, 1, sampling_prob)

  sim_2[i, ] <- x
  x[idx] <- sample

  idx <- ifelse(idx == 2, yes = 1, no = 2)
}

```

Subtask (f)


```

par(mfrow=c(2,1), mar = c(4,4,0,1))

plot(cumsum(sim_1[, 1])/c(1:50000),
     type = 'l', ylim = c(0, 1),
     ylab = "Frequency",
     xlab = "Iterations")
lines(cumsum(sim_2[, 1])/c(1:50000), col = "red")

plot(cumsum(sim_1[, 2])/c(1:50000),
     type = 'l', ylim = c(0, 1),
     ylab = "Frequency",
     xlab = "Iterations")
lines(cumsum(sim_2[, 2])/c(1:50000), col = "red")

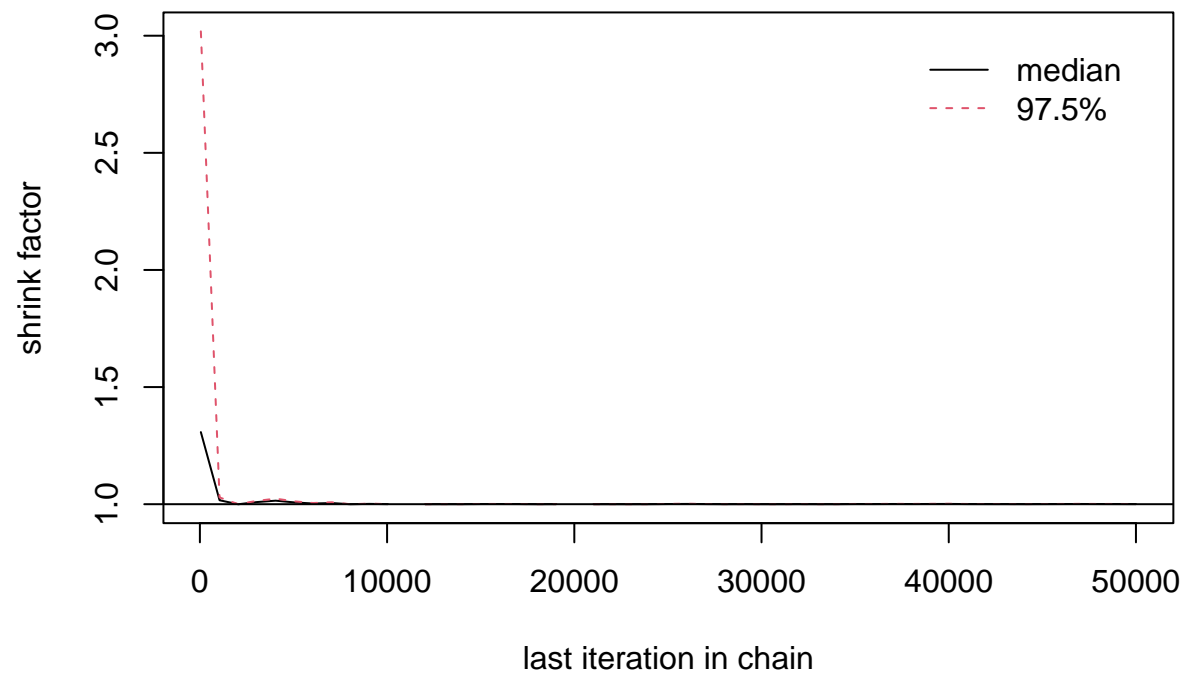
```




Based on the plots, we suggest a burn-in time of maximally a few thousand iterations. We pick 10000, as it is clear that this point is far past the burn-in phase, but implementing a threshold might allow for a more exact estimation.

Subtask (g)

```
gelman.plot(mcmc.list(mcmc(sim_1[, 1]), mcmc(sim_2[, 1])))
```



```
gelman.plot(mcmc.list(mcmc(sim_1[, 2]), mcmc(sim_2[, 2])))
```



We can see that the shrink factor converges to 1 fast after less than a few thousand iterations. This is in accordance with the result obtained in subtask (f).

Subtask (h)

This allows us to calculate the marginal probability $P(R = T | S = T, W = T)$:

```
# For simulation 1:
counts <- matrix(data = c(0, 0, 0, 0),
                 nrow = 2, ncol = 2)

for (i in 1:50000) {
  temp <- sim_1[i, ]
  temp <- temp + 1
  counts[temp[1], temp[2]] <- counts[temp[1], temp[2]] + 1
}

# Calculate probability:
(task_h_1 <- rowSums(counts/50000)[2])
```

```
## [1] 0.04936
```

```
# For simulation 2:
counts <- matrix(data = c(0, 0, 0, 0),
                 nrow = 2, ncol = 2)
```

```

for (i in 1:50000) {
  temp <- sim_2[i, ]
  temp <- temp + 1
  counts[temp[1], temp[2]] <- counts[temp[1], temp[2]] + 1
}

# Calculate probability:
(task_h_2 <- rowSums(counts/50000)[2])

```

```
## [1] 0.04628
```

Subtask (i)

Given the answers from subtask (a), the analytical solution returns:

$$\begin{aligned}
 P(R = T | S = T, W = T) &= \frac{P(R = T, S = T, W = T)}{P(S = T, W = T)} \\
 &= \frac{\sum_C P(C, R = T, S = T, W = T)}{\sum_{C,R} P(C, R, S = T, W = T)} \\
 &= \frac{\sum_C P(C)P(R = T|C)P(S = T|C)P(W = T|S, R)}{\sum_{C,R} P(C)P(R|C)P(S = T|C)P(W = T|S, R)}
 \end{aligned}$$

This returns:

```

# COmpute analytical solution
# Compare to:

task_c

```

```
## [1] 0.06
```

```
task_h_1
```

```
## [1] 0.04936
```

```
task_h_2
```

```
## [1] 0.04628
```